

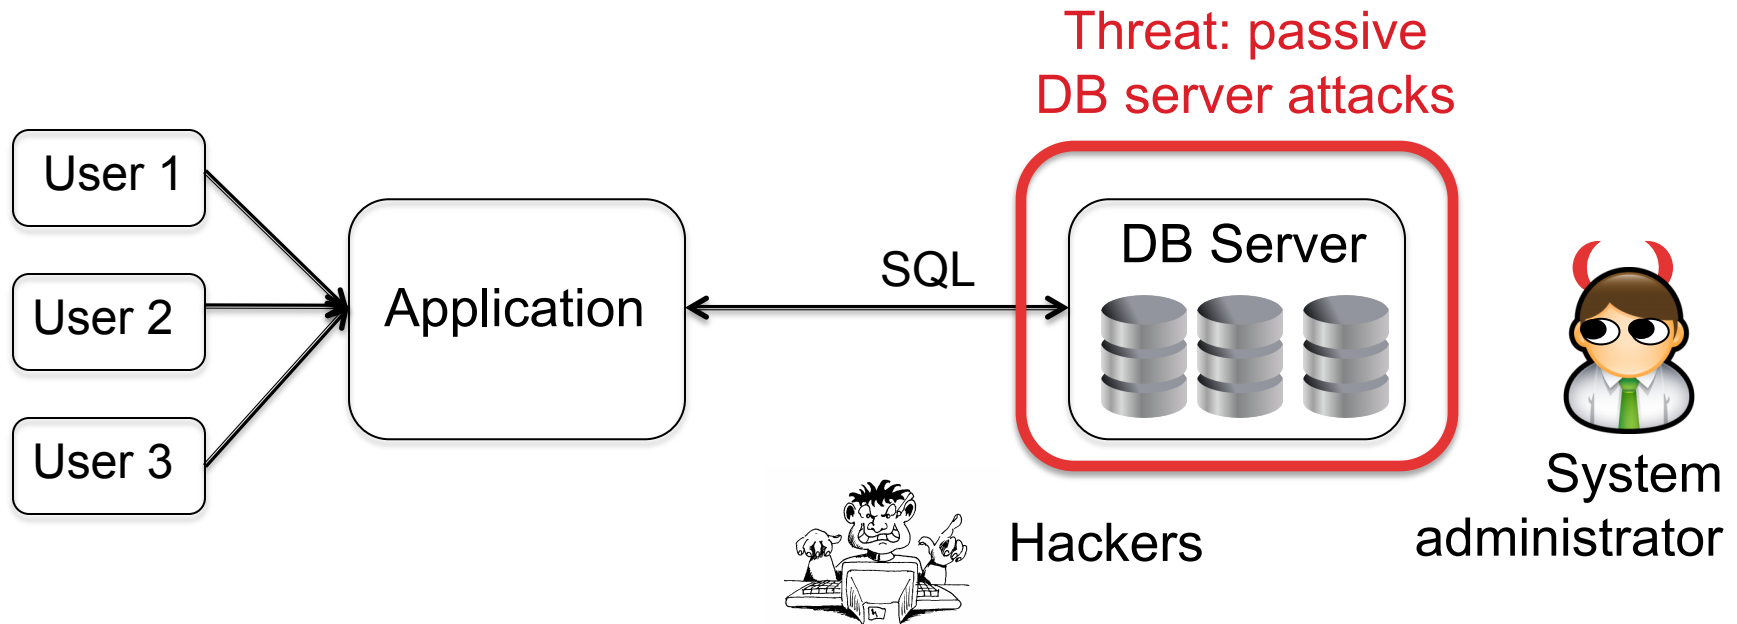
# CryptDB: Processing Queries on an Encrypted Database

Raluca Ada Popa, Catherine M. S. Redfield,  
Nickolai Zeldovich, and Hari Balakrishnan  
MIT CSAIL




# Problem

- ▶ Confidential data leaks from databases (DB)
  - ▶ 2012: hackers extracted 6.5 million hashed passwords from the DB of LinkedIn




- ▶ Process SQL queries on encrypted data

# Contributions

1. First *practical* DBMS to process most SQL queries on encrypted data  
 Hide DB from sys. admins., outsource DB to the cloud
2. Modest overhead: 26% throughput loss for TPC-C
3. No changes to DBMS (e.g., Postgres, MySQL) and no changes to applications

query  
x98aa = ?

salary
x4be2
x95c6
x2ea8
x17ce
...



Most SQL uses a limited set of operations

Security: Reveal only relations among data that are required by queries at column granularity

Unencrypted databases


CryptDB

FHE

query  
100 = ?


salary
60
100
800
100
...

fast  
insecure

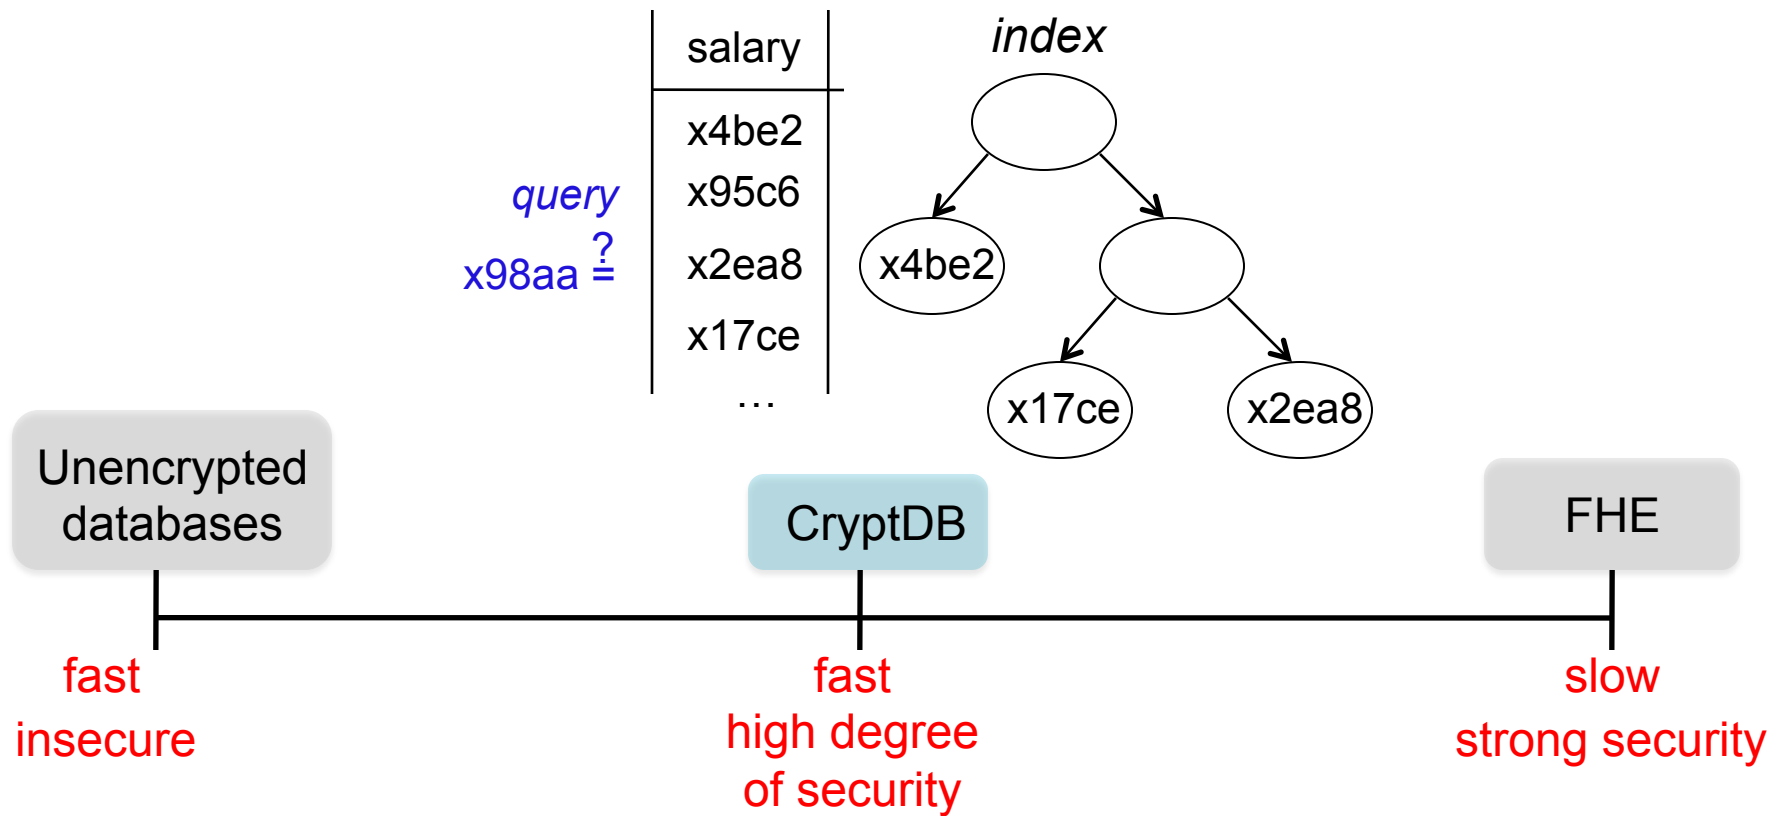


fast  
high degree  
of security

[Gentry'09], [GHS'12],...



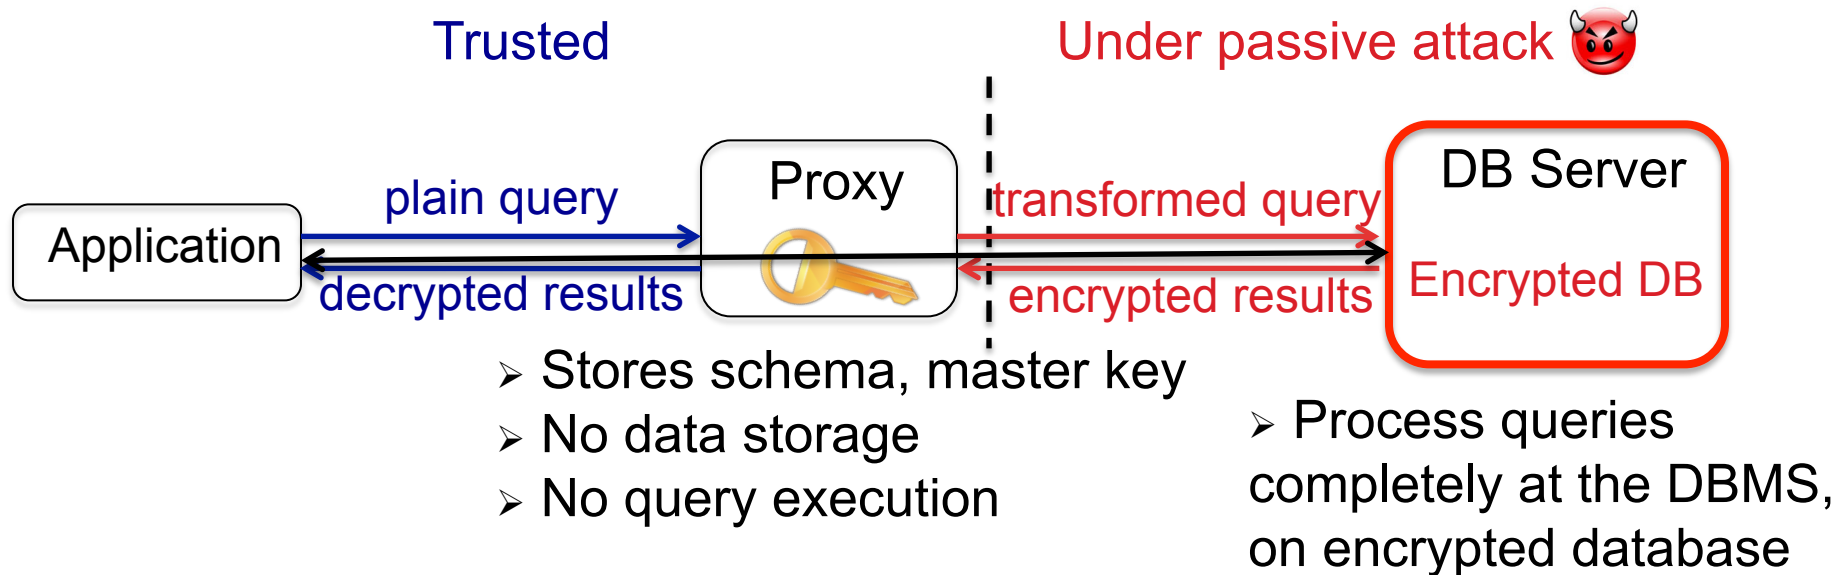
slow  
strong security



Other work: weaker security, functionality, and/or efficiency:

- Search on encrypted data (e.g., [Song et al., '00])
- Systems proposals (e.g., [Hacigumus et al., '02])
  - Require significant client-side processing

# System Setup



Application

SELECT \* FROM emp  
WHERE salary = 100

Proxy

SELECT \* FROM table1  
WHERE col3 = x5a8c34

table1/emp

col1/rank	col2/name	col3/salary
		x934bc1b
		x5a8c34
		x84a21c
		x5a8c34

~~Base table~~  
encryption

x5a8c34

x5a8c34

60

100

800

100

Application

SELECT \* FROM emp  
WHERE salary  $\geq$  100

Proxy

SELECT \* FROM table1  
WHERE col3  $\geq$  x638e54

~~Deterministic~~  
encryption

table1 (emp)

col1/rank	col2/name	col3/salary	
		x1eab81	60
		x638e54	100
		x922eb4	800
		x638e54	100

x638e54

x922eb4


x638e54



# Two techniques


1. Use SQL-aware set of encryption schemes
2. Adjust encryption of database based on queries

# Encryption schemes

<p>Highest</p>  <p>Security</p>	Scheme	Construction	Function	
	RND	AES in CBC	none	
	HOM	Paillier	+	e.g., sum
	SEARCH	Song et al., '00	word search	restricted ILIKE
	DET	AES in CMC	equality	e.g., =, !=, IN, COUNT, GROUP BY, DISTINCT
	JOIN	our new scheme	join	
	OPE	BCLO'09 + our new scheme	order	e.g., >, <, ORDER BY, SORT, MAX, MIN, GREATEST


# JOIN

- ▶ Do not know columns to be joined a priori!





- ▶ *KeyGen* (sec. param): SK
- ▶ *Encrypt* (SK, m, col i):  $C_m^i$  (with ) - deterministic
- ▶ *Token* (SK, col i, col j):  $(t_i, t_j)$
- ▶ *Adjust*  $(t_i, C_m^i)$ :  $C_m$  (with )

# JOIN (cont'd)



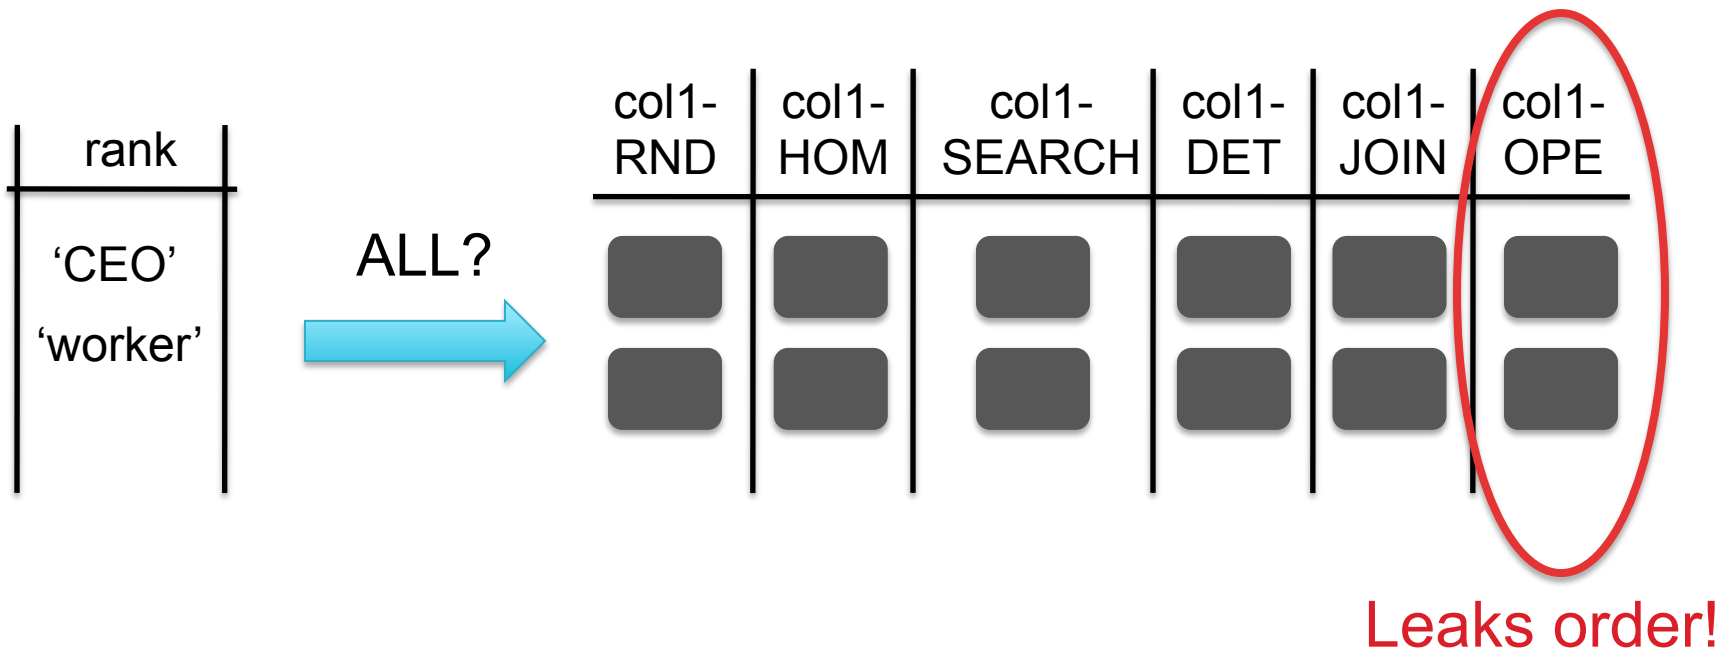
- ▶ Security: do not learn join relations without token
- ▶ Implementation:
  - ▶ 192 bits long, 0.52 ms encrypt, 0.56 ms adjust

# Encryption schemes

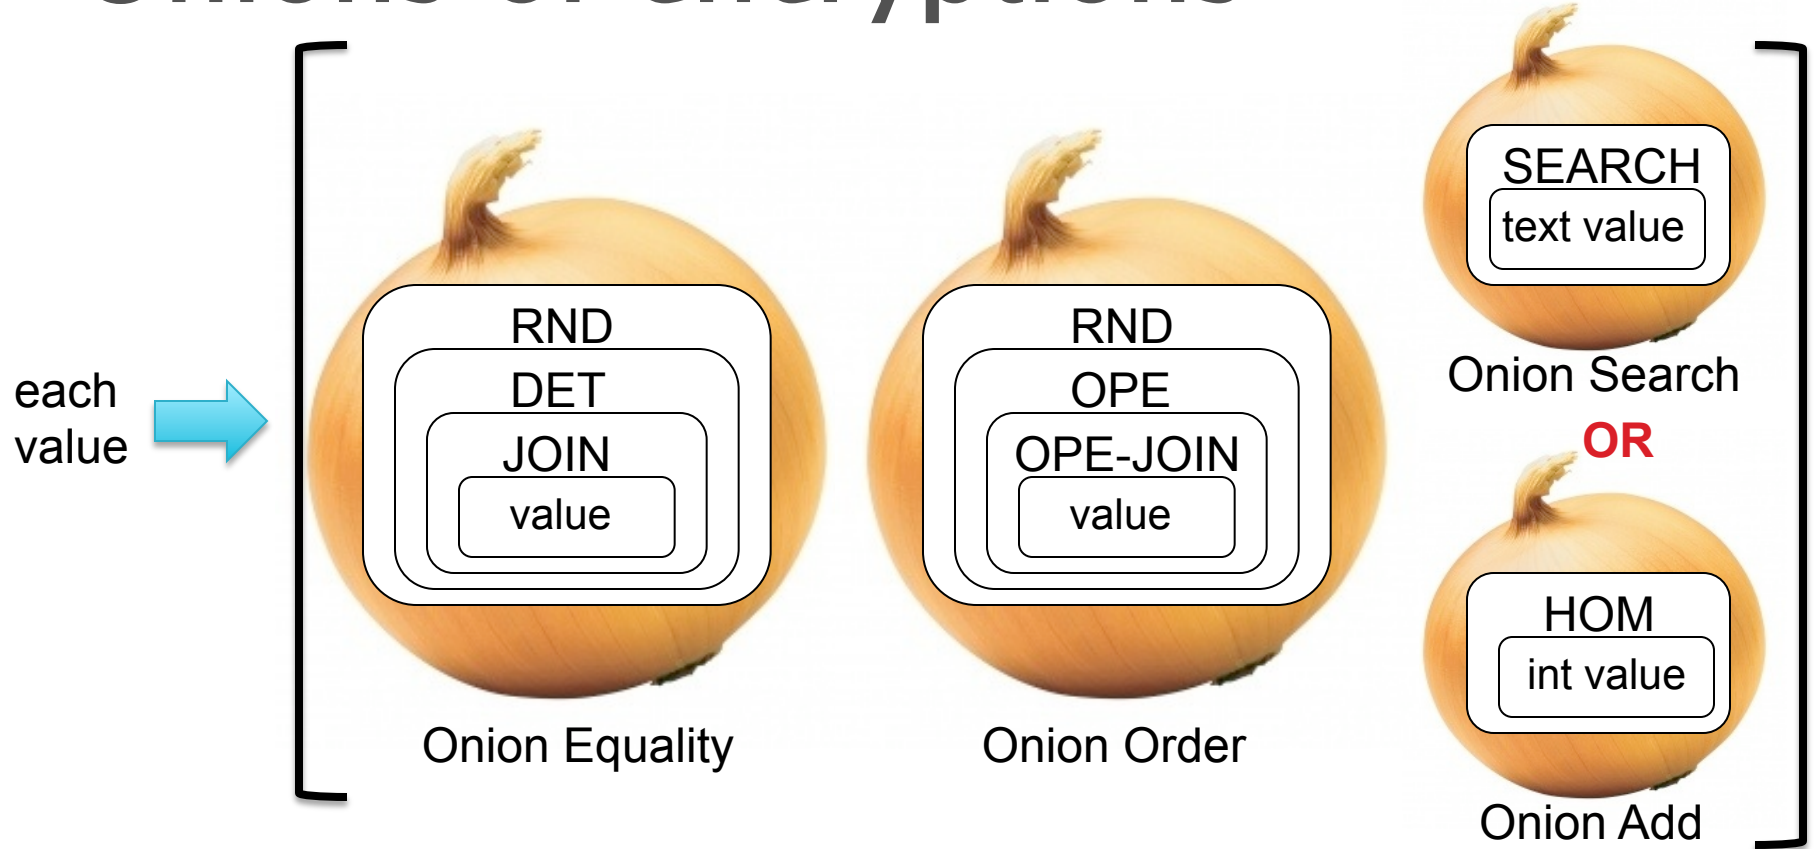
<p>Highest</p>  <p>Security</p>	Scheme	Construction	Function
	RND	AES in CBC	none
	HOM	Paillier	+, *
	SEARCH	Song et al., '00	word search
	DET	AES in CMC	equality
	JOIN	our new scheme	join
	OPE	Boldyreva et al. '09 + our new scheme	order
 <p>Functionality</p>			

# How to encrypt each data item?

- Encryption schemes needed depend on queries
- May not know queries ahead of time



# Onions of encryptions



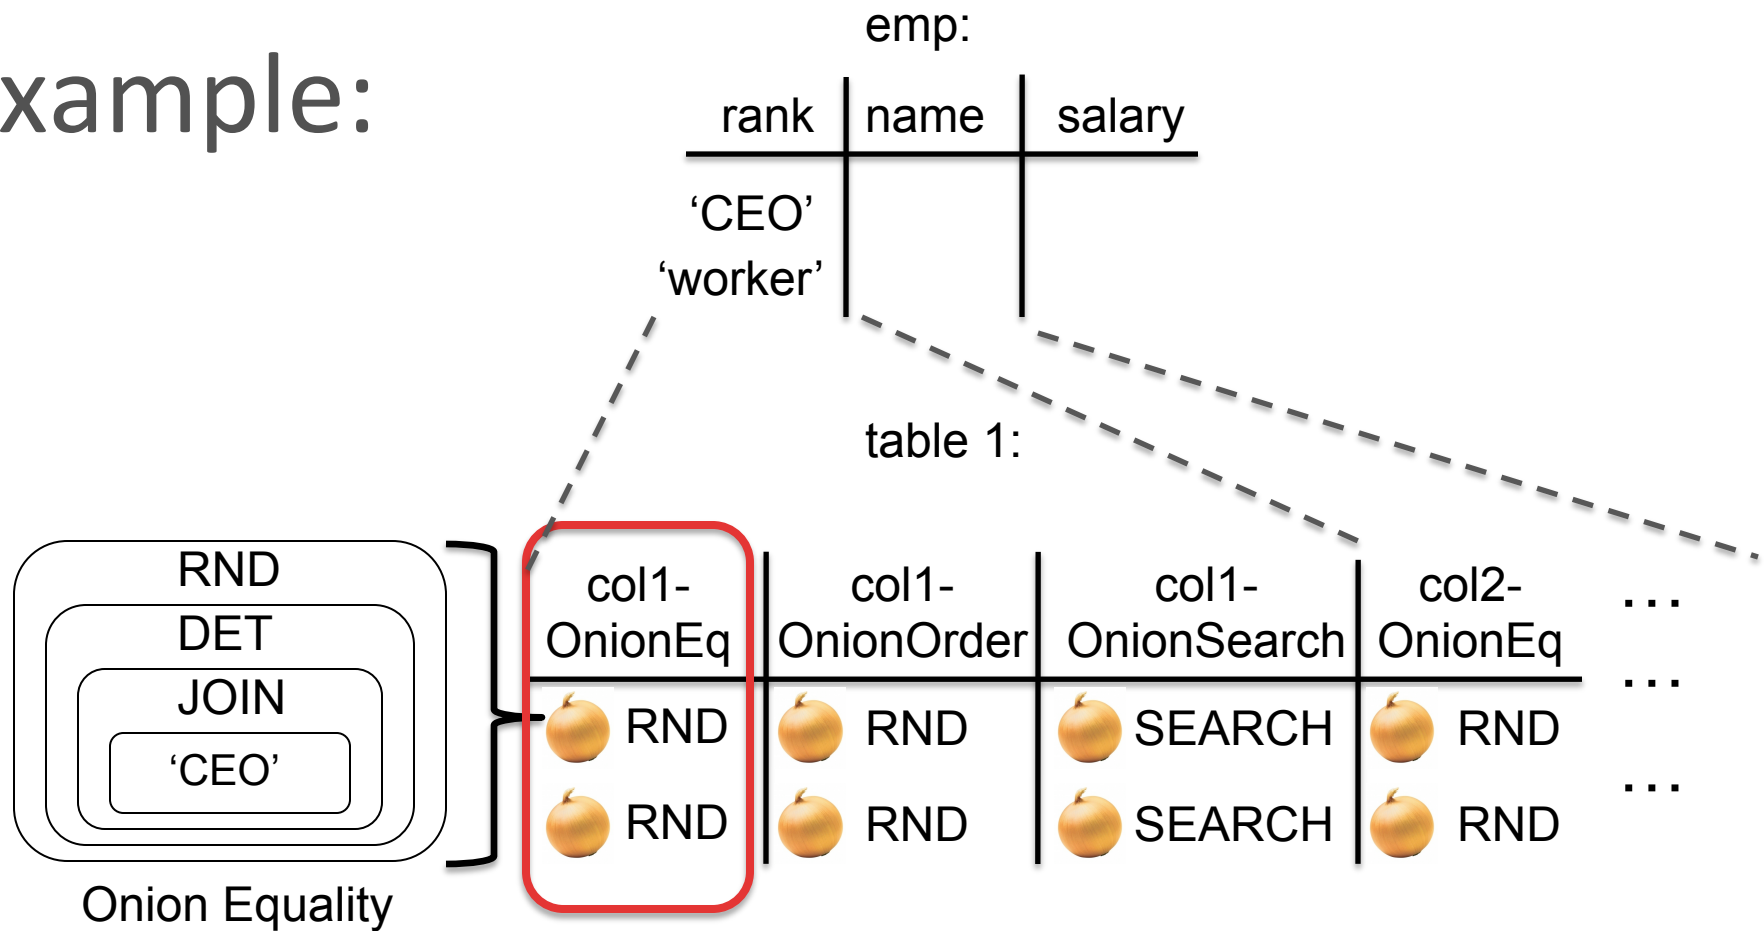
- Same key for all items in a column for same onion layer
- Start out the database with the most secure encryption scheme

# Adjust encryption

- Strip off layers of the onions
  - Proxy gives keys to server using a SQL UDF (“user-defined function”)
  - Proxy remembers onion layer for columns
- Do not put back onion layer




# Example:



SELECT \* FROM emp WHERE rank = 'CEO';

# Example (cont'd)



SELECT \* FROM emp WHERE rank = 'CEO';



UPDATE table1 SET col1-OnionEq =

**Decrypt\_RND**(key, col1-OnionEq);

SELECT \* FROM table1 WHERE col1-OnionEq = xda5c0407;

# Security guarantees

Queries → encryption schemes → leakage

➤ Encryption schemes exposed for each column are the most secure enabling queries

➤ Overall: Reveal only data relations needed for query type, at column granularity

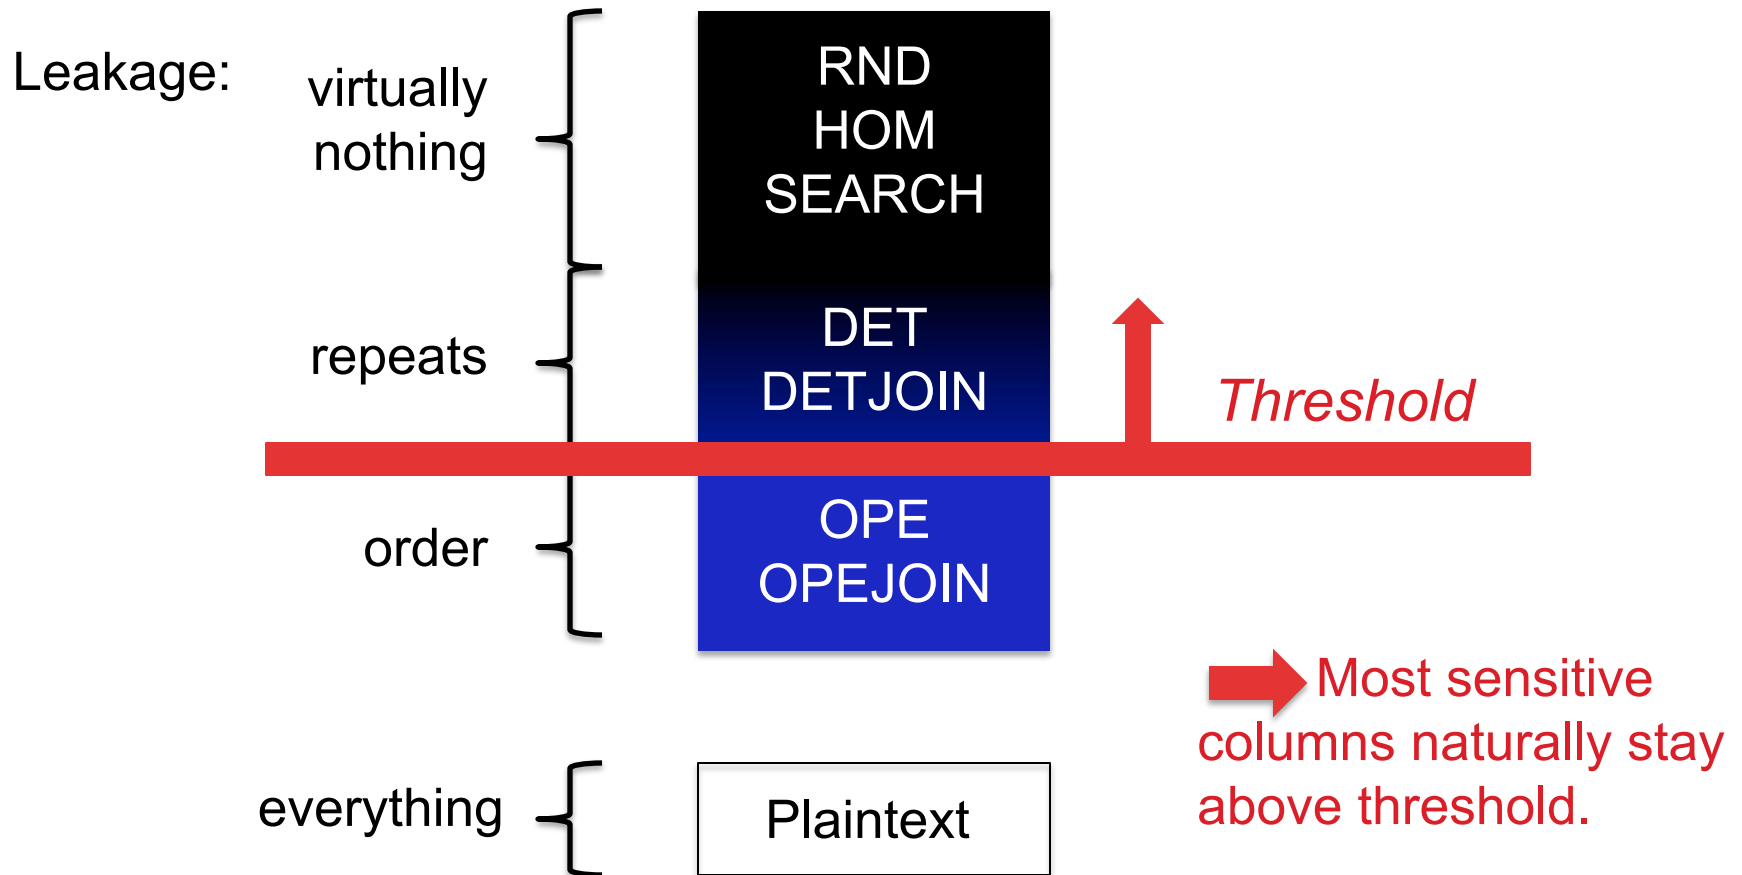
- equality predicate on a column → DET → repeats
- aggregation on a column → HOM → nothing
- no filter on a column → RND → nothing

*common in practice*

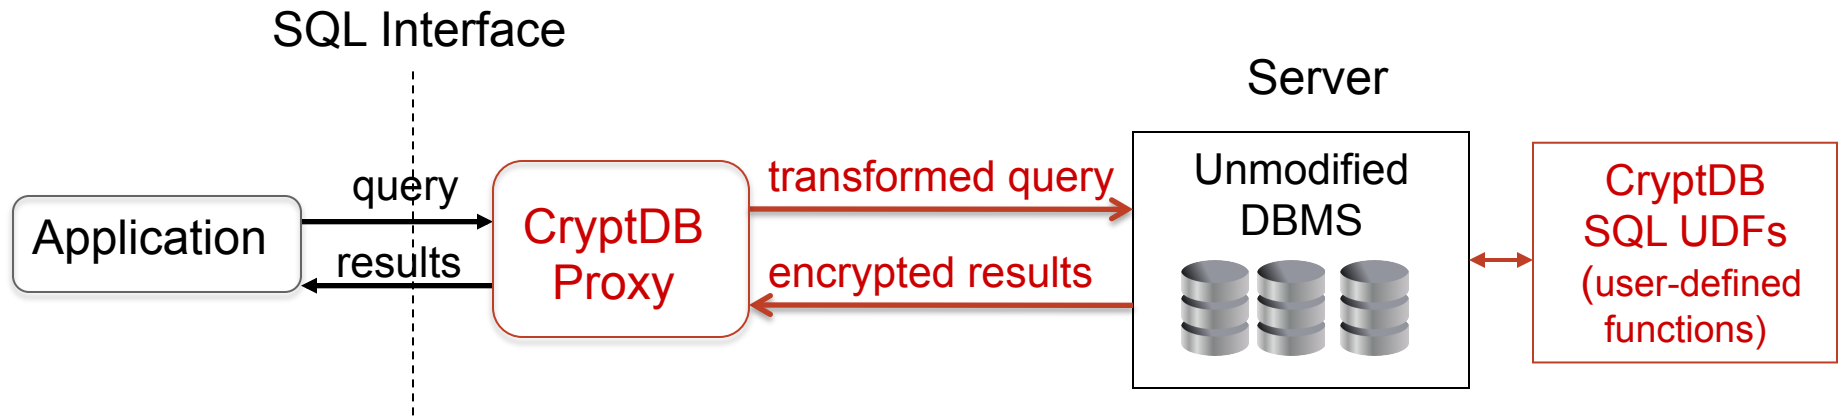
➡ Never reveals plaintext

# Security threshold

**SSN column**  $\geq$  *repeats*



# Implementation



- No change to the DBMS
- **Portable:** from Postgres to MySQL with 86 lines
- No change to applications

# Evaluation

1. Does it support real queries/applications?
2. What is the resulting confidentiality?
3. What is the performance overhead?

# Queries not supported

- More complex operators, e.g., trigonometry
- Operations that require combining encryption schemes
  - e.g.,  $T1.a + T1.b > T2.c$



**Extensions:** split queries, precompute columns, use FHE or other encryption schemes

# Real queries/applications

Application	Total columns	Encrypted columns	# cols not supported
phpBB	563	23	0
HotCRP	204	22	0
grad-apply	706	103	0
TPC-C	92	92	0
sql.mit.edu	128,840	128,840	1,094

SELECT 1/log(series\_no+1.2) ...  
... WHERE sin(latitude + PI()) ...



# Resulting confidentiality

Application	Total columns	Encrypted columns	Min level is RND	Min level is DET	Min level is OPE
phpBB	563	23	21	1	1
HotCRP	204	22	18	1	2
grad-apply	706	103	95	6	2
TPC-C	92	92	65	19	8
sql.mit.edu	128,840	128,840	80,053	34,212	13,131


Most columns at RND

Most columns at OPE analyzed were less sensitive

# Performance

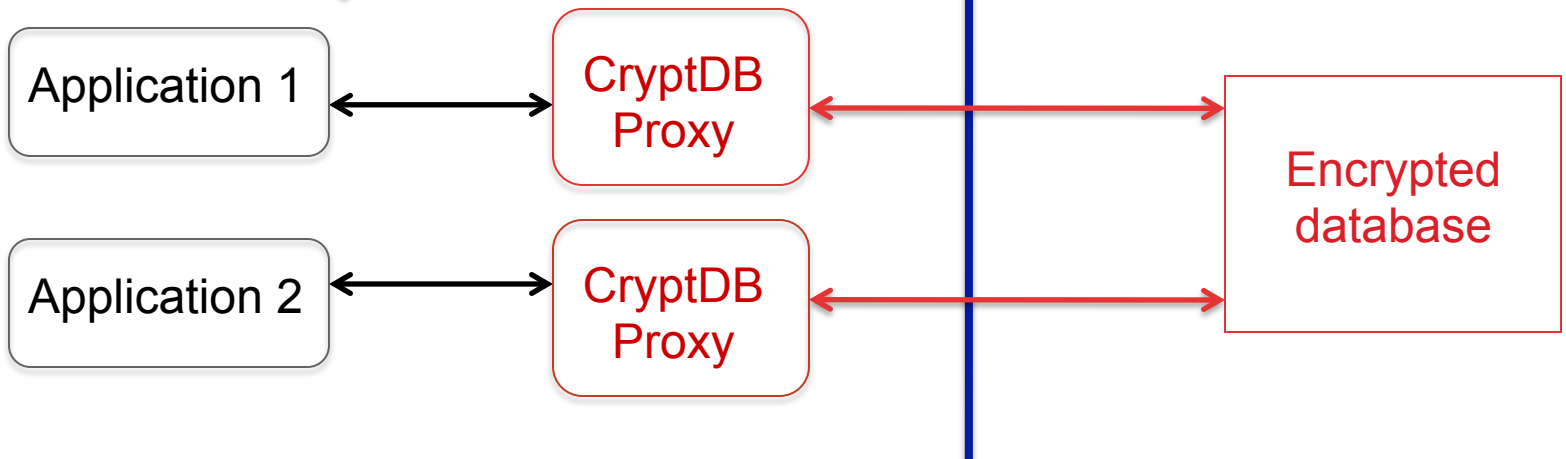
MySQL:

*DB server throughput*



*Latency*


CryptDB:



- Hardware: 2.4 GHz Intel Xeon E5620 – 8 cores, 12 GB RAM

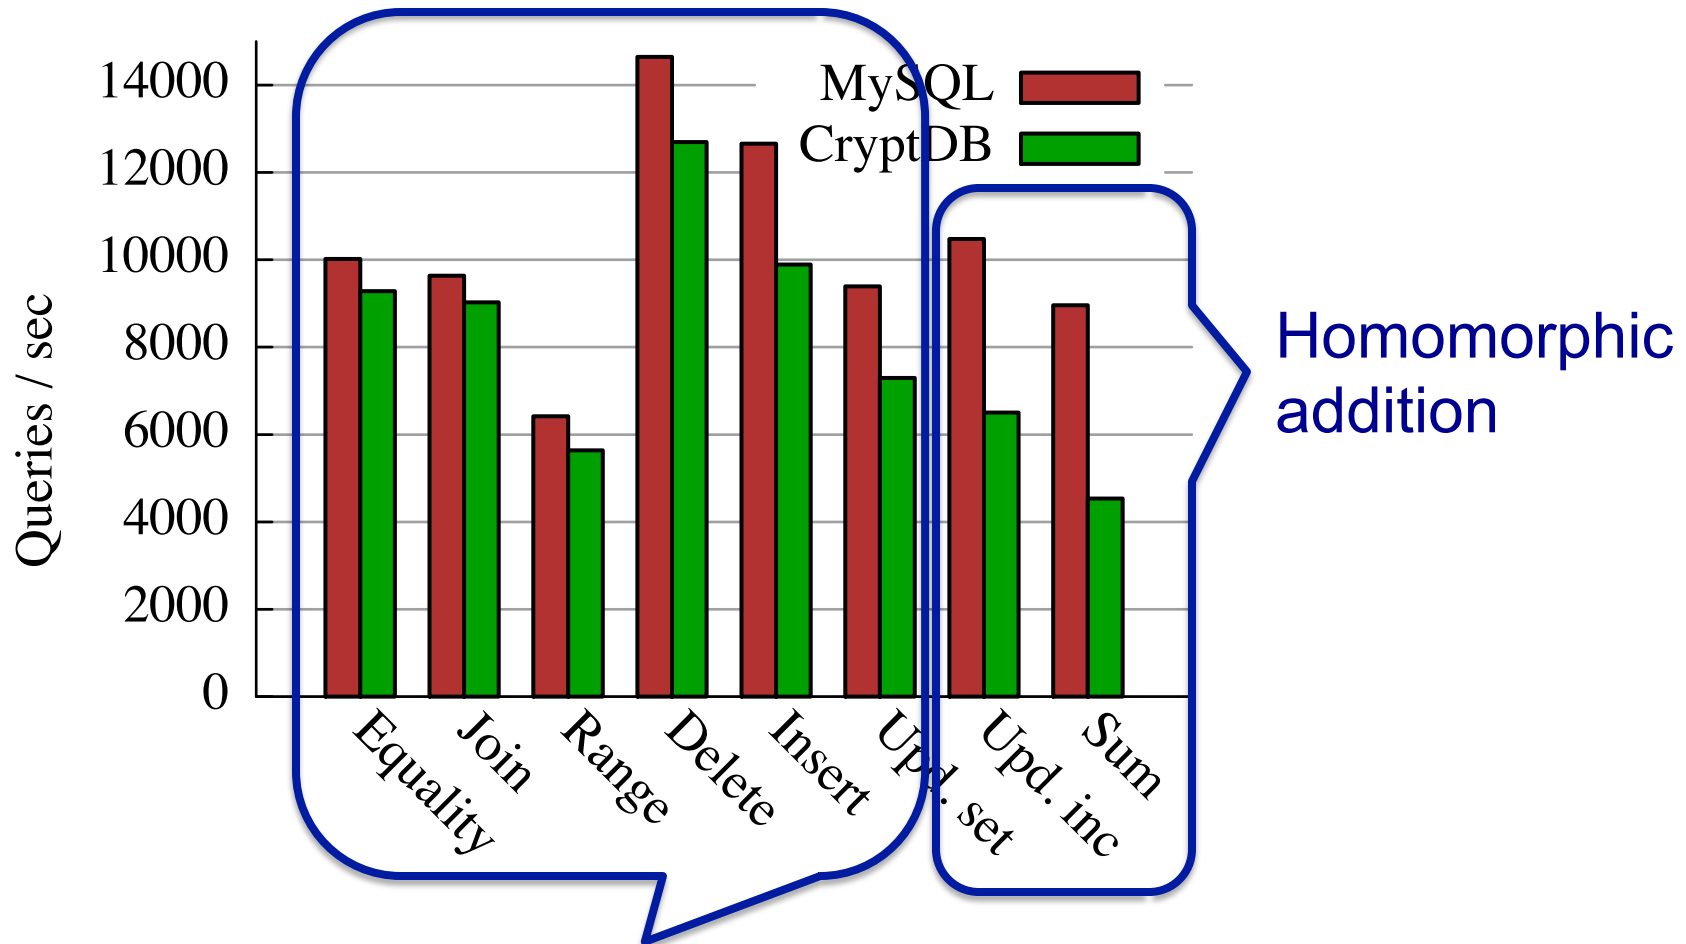
# TPC-C performance

➤ Latency (ms/query): 0.10 MySQL vs. 0.72 CryptDB



Throughput  
loss 26%

# TPC-C microbenchmarks



No cryptography at the DB server in the steady state!

➡ CryptDB is practical

Demo

# Conclusions

## CryptDB:

1. The first practical DBMS for running most standard queries on encrypted data
2. Modest overhead and no changes to DBMS

Website: <http://css.csail.mit.edu/cryptdb/>

# Thanks!